

Computer Science

Y. LIOW (MAY 7, 2024)

Contents

105 Containers and arrays	5001
105.1 Containers <small>debug: containers.tex</small>	5002
105.2 C++ integer types, size type, and midpoints <small>debug: cpp-integer-types.tex</small>	5004
105.2.1 Integer types	5004
105.2.2 Size type	5005
105.2.3 Midpoints	5007
105.3 Key and satellite data <small>debug: key-and-satellite-data.tex</small>	5010
105.4 Operations on containers <small>debug: operations-on-containers.tex</small>	5012
105.5 Static unsorted array <small>debug: static-unsorted-array.tex</small>	5013
105.5.1 Insert	5015
105.5.2 Delete	5018
105.5.3 Search	5019
105.5.4 Capacity issue	5021
105.6 Static sorted array <small>debug: static-sorted-array.tex</small>	5025
105.6.1 Insert	5025
105.6.2 Delete	5025
105.6.3 Search	5025
105.7 Dynamic array <small>debug: dynamic-array.tex</small>	5027
105.8 Memory fragmentation <small>debug: memory-fragmentation.tex</small>	5028
105.9 Accessing values in the container: index values <small>debug: index-values.tex</small>	5031
105.10 Accessing values in a container: pointers <small>debug: pointers.tex</small>	5032
105.11 Accessing values in a container: iterators <small>debug: iterators.tex</small>	5035
105.12 C++ STL: iterators and constant iterators <small>debug: constant-and-non-constant-iterator.tex</small>	5038
105.13 C++ STL: <code>std::vector</code> operations/methods using iterators <small>debug: vector-methods-with-iterators.tex</small>	5041
105.14 C++: Automatic type deduction <small>debug: type-deduction.tex</small>	5044
105.15 C++: Range-based loops <small>debug: range-based-for-loops.tex</small>	5045
105.16 C++ STL: <code>std::pair</code> and <code>std::tuple</code> <small>debug: products-and-tuples.tex</small>	5050
105.17 C++: typedefs <small>debug: typedefs.tex</small>	5052
105.18 C++: initializer lists <small>debug: initializer-lists.tex</small>	5054
105.19 C++ STL <code>std::vector</code> exercises <small>debug: exercises.tex</small>	5057

Chapter 105

Containers and arrays

105.1 Containers debug: containers.tex

From your programming experience, you see that your programming languages provide some basic types (`int`, `double`, `bool`, etc.) with operations/operators and functions. You then build a software using these types. On the way to building your software system, you will find that you usually have to build new types (such as structures) and functions to be used by your software. Some times, you will find that after some thought, certain functions can be placed inside structures, i.e., you combine data and functions, to get classes.

For instance if you want to develop an RSA crypto system, you will need to perform computations on integers with arbitrary number of digits. Well ... it cannot be arbitrary since you are limited by your computer system resources: you would need to work with integers with hundreds of digits. C/C++ does not allow you to do that. To model integers with arbitrary number of digits, you would probably use heap-based arrays, i.e., you would probably want to use pointers to `int` arrays in the free store. Since you expect to perform the usual arithmetic operations on such integers, you would probably want to develop a class called, say, `LongInt`.

Each `LongInt` object is essentially a container of digits. (Technically speaking it contains a pointer that points to a container – i.e., array – of digits in the heap. But you get the point.)

In your arithmetic operators (addition, subtraction, blah-blah-blah) you see that you frequently need to *travel through* the container, i.e., you will see something like this frequently in your code:

```
for (int i = 0; i < n; i++)
{
    ... do something with p[i] ...
}
```

if `p` is the pointer in your class and `n` is the number of digits `p` points to.

The technical CS term is that you have to **traverse** the container.

traverse

Since the heap-based array is linear, i.e., the layout of the data in the container is along a straight line, you see at least two different ways to travel through (traverse) the container. It's either

```
for (int i = 0; i < n; i++) // forward traversal
{
    ... do something with p[i] ...
}
```

or

```
for (int i = n - 1; i >= 0; i--) // reverse traversal
{
    ... do something with p[i] ...
}
```

So in this case, there are two obvious traversals.

You can also fantasize about something like this:

```
for (int i = 0; i < n; i += 2) // forward, even index
{
    ... do something with p[i] ...
}
```

Anyhow, you see that your `LongInt` class will need to have a container and you need to traverse the container.

Other things you want to do with a container is to put things into the container, remove things from the container, search for a value in the container, etc.

If the above idea of a container applies only to your `LongInt` class, then there's no reason to make a big deal out of it and create a whole theory or body of knowledge of such thingies. But it turns out that containers of different shapes (i.e. not necessarily along a straight line) is actually very common in nature (i.e., in computations).

When I say that containers are common in nature I really mean it. Think of your bookshelves. As a whole, it's a container of books. When you need a book, you have a search for it and then remove it from your shelves. Of course you probably want to add books to your shelves too.

In the same way, a database is a container of data. For scientific purposes, it could be a database of genomic data; for business applications, it could be a database of customer, product, and sales transaction data. You also want to add data, remove data, and search for data in databases.

In a computer game, you might have a container of all spaceships attacking you. In a fantasy role playing game, you might have an inventory list (weapons and what-nots) which is a container.

So I'm not lying nor exaggerating when I say that the world is full of containers.

By the way, although beginning algorithms courses focus on the algorithms operating on containers, you should know that the study of algorithms is broader than that.

105.2 C++ integer types, size type, and midpoints

debug: cpp-integer-types.tex

105.2.1 Integer types

You should already know (from CISS240, CISS245) that there are several integer types.

- `int`
- `long int`
- `long long int`
- `unsigned int`
- `unsigned long int`
- `unsigned long long int`

`char` is actually also an integer type, which is why `char` values can be case labels. There's also `unsigned char` type. `bool` is also an integer type.

The problem with the `int`, `long int`, `long long int` is that the number of bits used by them is not standardized. For me (C++ code compiled with g++ on F31 vm), `int` is 32 bits and an `int` value lies in the range $-2^{31}, \dots, 2^{31} - 1$ and `unsigned int` is 32 bits and an `unsigned int` value lies in the range $0, \dots, 2^{32} - 1$. On such systems, the `long long int` is 64 bits and the `long int` can be either 32 bits or 64 bits. In the future, an `int` might be 64 bits wide. (When I was in high school, my PC had an Intel 80286 CPU which had a 16-bit architecture, i.e., an `int` then had 16 bits.) This can be a problem: if you bring your C++ code from a machine where an `int` is 32 bits in length and recompile it on another machine where the `int` is 16 bits in length, your program might not work correctly when it processes an integer such as 2^{16} . (Why?)

To make your program portable (i.e., can be compiled and run on any machine that has a C/C++ compiler), C/C++ provides integer types with fixed bit lengths.

- `int8_t`, `uint8_t`
- `int16_t`, `uint16_t`
- `int32_t`, `uint32_t`
- `int64_t`, `uint64_t`

The “`_t`” in the above means “type”. There are constants for the limits of the above types. For instance:

- For `int32_t`, the minimum is `INT32_MIN` and the maximum is `INT32_MAX`.
- For `uint32_t`, the minimum is 0 and the maximum is `UINT32_MAX`.

Professionally written C++ code tends to use these types. If you want to use the above, you need to “`#include <stdint>`”.

105.2.2 Size type

Another thing to know is that for containers (including `std::vector`), we have the concept of the size of a container. Recall that if `v` is a `std::vector< int >` object, the size of `v` is

`v.size()`

The type of `v.size()` and `v.capacity()` is `size_t` (“size type”). `size_t` is an unsigned integer type and is either `uint32_t` or `uint64_t`, depending on the platform. This makes sense since you don’t expect a size to be negative!!! (It doesn’t make sense to say a vector has -5 values.)

When you write something like this:

```
for (int i = 0; i < v.size(); ++i)
{
    ...
}
```

you’ll get a warning since `i` is an `int` and `v.size()` is an unsigned `int`:

```
main.cpp:5:23: warning: comparison of integer expressions of different
signedness: ‘int’ and ‘std::vector<int>::size_type’ {aka ‘long unsigned
int’} [-Wsign-compare]
    5 |         for (int i = 0; i < v.size(); ++i)
      |           ~~~~~~
```

To be absolutely correct, one would write

```
for (unsigned int i = 0; i < v.size(); ++i)
{
    ...
}
```

or even better

```
for (size_t i = 0; i < v.size(); ++i)
{
    ...
}
```

If you want to use `size_t` you might have to do “`#include <cstddef>`”:


```
#include <cstdint> // remove this and you'll get an error
int main()
{
    size_t i = 0;
    return 0;
}
```

However if you are already using any C++ STL container classes such as `std::vector`, the container class has already done “`#include <stdint>`” since `std::vector` class uses `size_t`.

In some cases, it won't make a difference whether you are using `int` values or `unsigned int` values. In particular for 32-bit `int` and 32-bit `unsigned int`, their overlap is $0, \dots, 2^{31} - 1$. Because of the math of binary representations of integers and unsigned integers (see CISS360), it doesn't matter if you are using `int` or `unsigned int` if you stay in $0, \dots, 2^{31} - 1$. So it's perfectly OK to write

```
for (int i = 0; i < v.size(); ++i)
{
    std::cout << x[i] << '\n';
}
```

if your computations involving `i` and `v.size()` stays in $0, \dots, 2^{31} - 1$, except that you might get compiler warnings.

One final point about working with size type and type for index values. Recall from CISS240 when we use `int` for index type, we frequently use `-1` to indicate “invalid index value”. For instance during a search (linear search or binary search) when a target value is not found in an array, we return `-1`. What will happen if you want to handle such scenarios and you are using `size_t` or `unsigned int` for your index type?

When I run this:

```
#include <iostream>

int main()
{
    std::cout << sizeof(size_t) << '\n';
    size_t x = -1;
    std::cout << x << '\n';
    return 0;
}
```

I get

```
8
18446744073709551615
```

(This is on our f31 vm). What's happening? The above tells me that a `size_t` value consumes 8 bytes, i.e., 64 bits. That huge positive integer 18446744073709551615 is of course then $2^{64}-1$. (Check this yourself). The `-1` (as an `int`) value is the same as 18446744073709551615 (as a `size_t` value). This is due to “clock arithmetic”, the way math works in our CPU. I already talked about this in CISS240 when I talked about the “clock arithmetic” of the `int` type. And “clock arithmetic” is the informal term for “modular arithmetic” (discrete math, CISS360, CISS451) which is the mathematical model used in designing CPUs.

This means that if you have a really huge array `x` and `x[18446744073709551615]` is actually a value in the array and some search algorithm returns `-1` to a `size_t` variable, you'll have a problem: does the algorithm mean that value is not found or does it mean the value is found at index 18446744073709551615?

These are things (among others) that a serious computer scientist have to be aware of.

By the way since the C++ string class has a definition of the constant `std::string::npos` which is `size_t (-1)`. The name of this constant `npos` reads “not a position”, “no position”, “non-position”.

105.2.3 Midpoints

Frequently you need to compute the midpoint of two indices. We saw that in binary search in CISS240. You will see a lot more of that. Given two index values say `left` and `right`, the midpoint index value can be computed as

$$\text{mid} = (\text{left} + \text{right}) / 2$$

However there's problem with this if `left` and `right` are huge – `left + right` will give you an arithmetic overflow. If `left` and `right` are unsigned int of 32-bits, then `left + right` overflows if `left` and `right` are both slightly smaller than $2^{32} - 1$. A better mid point calculation is

$$\text{mid} = \text{left} + (\text{right} - \text{left}) / 2$$

Make sure you run this program:

```
#include <iostream>

int main()
{
    unsigned int left = 0 - 3; // 2^32 - 3
    unsigned int right = 0 - 1; // 2^32 - 1
    unsigned int mid = 0 - 2; // correct
    std::cout << left << ' ' << mid << ' ' << right << '\n';
    mid = (left + right) / 2; // wrong
    std::cout << left << ' ' << mid << ' ' << right << '\n';
    mid = left + (right - left) / 2; // correct
    std::cout << left << ' ' << mid << ' ' << right << '\n';
    return 0;
}
```

We also need to compute the midpoint between two addresses. Here's how you do it (similar to the above):

```
#include <iostream>

int main()
{
    int x[8] = {2, 3, 5, 7};
    int * left = &x[0];
    int * right = &x[4];
    int * mid = left + (right - left) / 2; // ERROR: (left + right) / 2 !!!
    std::cout << sizeof(int) << '\n';
    std::cout << left << ' ' << mid << ' ' << right << '\n';
    std::cout << (unsigned long long) left << ' '
                << (unsigned long long) mid << ' '
                << (unsigned long long) right << '\n';
    std::cout << *left << ' ' << *mid << ' ' << *right << '\n';
    return 0;
}
```

Note that for pointers, you cannot compute “(left + right) / 2” because you cannot add addresses. Try it out yourself and read the error message.

In case you want to store the difference of addresses, the type for difference of addresses is `std::ptrdiff_t` and if you want to use it do “`#include <cstdint>`”:

```
#include <iostream>
#include <cstdint>

int main()
{
    int x[8] = {2, 3, 5, 7};
    int * left = &x[0];
    int * right = &x[4];
    std::ptrdiff_t diff = right - left;
    return 0;
}
```

105.3 Key and satellite data debug: key-and-satellite-data.tex

In the case of an array, such as an array of integers representing the digits of a long integer, the index provides a search “key”. For instance “the third digit of the long integer represented by x ” (i.e., the hundreds) means $x[2]$. In this case “third” means “2” which leads us to $x[2]$.

In many applications, the key can be different (i.e., not an index value). For instance suppose I have array of students:

```
#include <iostream>
#include <string>

class Student
{
public:
    Student()
    {}

    Student(const std::string & id,
            const std::string & firstname,
            const std::string & lastname,
            : id_(id),
            firstname_(firstname),
            lastname_(lastname)
    {}

private:
    std::string id_;           // ----> key

    std::string firstname_;    // \
    std::string lastname_;     // +--> satellite data
    // dob, address, etc.      // /
};

int main()
{
    const int CAPACITY = 1000;
    Student student[CAPACITY];
    int n = 0;

    // Set student[0], student[1], student[2], ...,
    // student[94].
    // Set n to 95
    // Altogether we have 95 students in the system.

    return 0;
}
```

In this case, search for a particular student might depend on supplying the system with a student id. The value of the student id (for each object in the array) is stored in instance variable `id_`.

In this case student id is the key. A field that uniquely identifies a value in a container of values is called **key**. The other data which is not used to uniquely identify a particular entry in the array is called **satellite** data. In the above, the first name and last name for each entry in the array is satellite data of that entry. The value of `firstname` and `lastname` does not uniquely identify a student in the array since it's possible to have two students named John Doe. Right? In a real information system of students, one would expect a lot more satellite data than key data.

key
satellite

The key problem is how to design containers so that when given a key and satellite, you can quickly put the (key, satellite) into the container, when given a key, how to find the (key, satellite) in the container, and when given a key, how to delete the (key, satellite) from the container (if it's found at all). In particular, we will focus mainly on the organization of the keys in the container so that the operations are as fast as possible.

For large scale data, it's common to store the satellite data separate from the key data. In that case a **Student** object might have key data and some kind of data that leads to its satellite data. Why? Because then a huge amount of key data can be loaded into memory from harddrive so that search can be as fast as possible. Once the student's id is found, then only this student's satellite data is read from harddrive. Situations like this do occur. See for instance CISS430 (database).

105.4 Operations on containers debug: operations-on-containers.tex

What are some of the common operations you want to have on a container, such as for instance an array of students where the key is the student's id and the satellite data include the student's first name, last name, etc.?

- You want to **add** a value (example: a student object) to a container (example: a container of students).
- You want to **delete** a value from the container.
- You want to **search** for a student in the container.

The above are probably the most common and most important. Here are some more:

- When the container has a sense of ordering of values, you might want to do this: find the k -th smallest value or k -largest value in the container.
- When you have two containers of the same kind, you might want to merge both into one.

Etc.

In the case of an array, there *is* a concept of ordering. You have the concept of the “third student”, the “105-th student”, etc. If you prefer to call students the “zeroth” student, the “first” student, the “second” student, the “third” student, etc then the “third” student in the container is `student[3]`. Of course what's meant by “third” depends on how you order the values. In the case where the order is

`student[0], student[1], student[2], ..., student[n - 1]`

the k -th value is `student[k]`. However if I prefer to view the values like this (i.e., in reverse direction):

`student[n - 1], student[n - 2], student[n - 3], ..., student[0]`

then the k value means `student[n - k]`.

There are many, many, many other possible operations on the container, depending on the “structure” of the contains.

For instance if the container is a graph, i.e., think of this as a bunch of dots and lines where data is stored at the dots. Then one operation is “find the shortest path from one dot to another.”

105.5 Static unsorted array debug: static-unsorted-array.tex

I will say “static array” to mean arrays with memory allocated in the frame, i.e. not in the heap, because the amount of memory used for such things are fixed while your array is in scope.

Here’s an example:

```
const int N = 1000;
int x[N];
```

If T is a type (int, double, bool, structure, class, ...), you do the same thing:

```
const int N = 1000;
T x[N];
```

The size of the array (see N) is usually a constant although some C/C++ compilers allow you to use a variable. But in any case, after the array is requested, you cannot change the size of your x:

```
const int N = 1000;
T x[N];
...
gimme_more(x, 1000000); // NOPE!!!
```

This is probably the simplest container. Here are some of the things you can do with your array.

If you want to simulate a container containing different number of items in the container, you have to include a variable that records the number of things already placed in the container:


```
const int N = 1000;
T x[N];
int n = 0; // number of things in x, i.e., only x[0],...,
           // x[n-1] are considered values in the
           // container. The values x[n],...,x[N-1]
           // should be considered extra unused space.
```

When you package that into a class, it would look like this:



```
template < typename T >
class Array
{
private:
    static const int N = 1000;
    T x[N];
    int n;
};
```

105.5.1 Insert


If you need to insert a value at index position i , then you have to move the values at positions $i, \dots, n - 1$ to positions $i + 1, \dots, n$ to make an “empty space” for the new value. If i is 0, that requires moving lots of values. If i is $n - 1$, then only one value has to be moved. If i is n , the value is to be placed just one step past the last value in the array – you don’t have to move anything. Here’s the picture of insert at index 2 where n is 5 and the capacity (size) is 10:



The values from index values 2 to 4 are moved (i.e., copied) to the right by one step:



and then 8 (the new value) is placed at index 2:



Of course after the above is done, you have to increment n , the count of number of things in your array.

```
for (int j = n - 1; j >= i; j--)
{
    x[j + 1] = x[j];
}
x[i] = newvalue;
n++;
```


The time taken is

$$A + B(n - i)$$


for constants A and B .

Let me call the first value of the array the **head**. If a value is to be inserted at


index 0, I will say that I'm "inserting at the head" (to be precise, I'm inserting before the current head – the new value becomes the head of the new array). Pictorially, this is what happens: I have an array of values 2, 6, 4, 5, 3 and I want to insert an 8 at the head:




After doing that I get this:



Let me call the last value of the array the **tail**, i.e. at index $n - 1$. If a value is to be inserted at index $n - 1$, I will say that I'm "inserting a tail" (to be precise, I mean inserting after the current tail – it becomes the tail of the new array). Pictorially, this is what happens: I have an array of values 4, 2, 3, 1, 5 and I want to insert an 8 at the head:



After that I get



Inserting at the head takes the most time. Big-O-wise, insert at the head has a runtime of

$$O(n)$$

Inserting at the tail is the fastest since the runtime is

$$O(1)$$

Averaging over all possible $i = 0, 1, \dots, n - 1$ and $i = n$ (the case where you're

inserting just beyond the array) the average runtime is

$$\begin{aligned} & \frac{1}{n+1} ((A + Bn) + (A + B(n-1)) + \cdots + (A + B) + A) \\ &= \frac{1}{n+1} (A(n+1) + B(n + \cdots + 1)) \\ &= \frac{1}{n+1} \left(A(n+1) + B \frac{n(n+1)}{2} \right) \\ &= A + \frac{B}{2}n \\ &= \frac{B}{2}n + A \end{aligned}$$

I've used the formula:

$$1 + 2 + \cdots + n = \frac{n(n+1)}{2}$$

OK ... now we know that the average time taken to insert is $O(n)$. (Don't recall the big-O notation and all that algorithmic analysis goodies? Don't panic. Just check the relevant notes/chapters.) Of course I'm assuming in the above that all index values are equally likely.

Here are runtime complexities for insert:

- Insert head: $O(n)$
- Insert tail: $O(1)$
- Average insert: $O(n)$

(assuming the array is not full.)

105.5.2 Delete

If you want to remove a value, say it's at index position i , you have to move the value at $i + 1, \dots, n - 1$ to the “left” by one position to index positions $i, \dots, n - 2$. After that's done, you have to decrement your n . The average and worse runtime is again $O(n)$. The best case runtime is $O(1)$.

Here are the runtime complexities for delete:

- Delete head: $O(n)$
- Delete tail: $O(1)$
- Average delete: $O(n)$

105.5.3 Search

What about search (or find)? In the worse case, you have to scan the whole array. Which means that the time taken is $O(n)$. The average case is also $O(n)$. The best case runtime is $O(1)$.

Now suppose we sort the values in $x[0], \dots, x[n - 1]$ (say in ascending order). Depending on the sorting algorithm used (see relevant notes) the best time for sorting is $O(n \lg n)$. The benefit of having a sorted array is that search time using binary search takes

$$O(\lg n)$$

(see chapter on algorithmic analysis) which is better than $O(n)$. If you want to ensure that the array is always sorted, when you insert a value in the array, you (of course) don't have to specify where to insert the value but make sure that the value is inserted so that the list remains sorted. How much time do you need? Well ... first you do a binary search. If it's found, you can just insert at that place. You would have a duplicate value - if that's allowed in your list. If it's not found, a -1 is returned. Now you call a version of the binary search that returns the index of where a new value should go instead of -1 . (This modified binary search is not difficult to write.) You performed two binary searches. So the time taken is $O(2 \lg n) = O(\lg n)$. Then you perform the insert, taking $O(n)$ time. So all in all, insert also takes $O(n)$ time to insert. Of course instead of performing two binary searches, you can just do a linear search and perform the insert. All in all, you would still get $O(n)$ for runtime.

For delete, say you specify the value to delete and not the index. In this case, you again need to search for the index where the value occurs (taking $O(\lg n)$ time using binary search).

Unsorted array	Best	Average	Worse
Insert	$O(1)$	$O(n)$	$O(n)$
Delete	$O(1)$	$O(n)$	$O(n)$
Search	$O(1)$	$O(n)$	$O(n)$

Note that in many cases the best case runtime is not that useful. The average and worse is more important. When you build a system, you want to focus on the worse case scenario, right? The best case scenario probably does not occur frequently and it would be overly optimistic to focus only on that.

But what if we're in a situation where the best case always occurs? Maybe the program you're trying to build requires only insert and delete at one end of the array. In the case of the (unsorted) array, the best case for insert and delete

is pretty fast if you insert and delete at the tail end of the array. Remember that!!! (See a later chapter on stacks, queues, etc. implemented using arrays.)

There's one thing to note about the array: You can very quickly get to the i -th value in the array. The time taken is the same (pretty much) regardless of the value of i . You will see later that there are containers where going to the i -th value of the container is pretty costly. So this is a plus point for arrays.

105.5.4 Capacity issue

The *BAD* thing about the static array (and you absolutely have to remember this) is that the maximum size (or capacity if you like) of the container is fixed and cannot be changed during runtime. If I declare an array with $N = 1000$ values and I keep adding things to this array, ultimately my n will reach 1000 and I won't be able to add more stuff into the array. You might think ... "Well ... I'm going to start with 100000 then!" True, true, true. But what if 1000000000 things were intended to be put into the array? Another thing is this: What if there's a scenario where the number of things needed is actually 100? You would be wasting lots of memory. It's just difficult to write flexible software to handle different situations if you're using static arrays. So the static'ness of the memory usage of static arrays is a very serious disadvantage.

Exercise 105.5.1. Write a template `Array` class with the following (public! ... right?) methods:

```
void    insert(int index, T key);
void    remove(int index);
int     find(const T & key); // returns index or -1 (if not found)
const T & operator[](int index) const;
T & operator[](int index);
```

Allow duplicate values in an `Array` object. Also, write a template function to print the contents of your object:

```
template < class T >
std::ostream & operator<<(std::ostream & cout,  const Array &);
```

If the `Array< int >` object contains 1, 2, 3, then printing it will give you

```
[1, 2, 3]
```

Set the maximum capacity of your `Array` objects to 1000000. Remember my advice in class: Always write a non-template version first!!! Collect timing data for the following scenarios for `Array< int >` objects:

- What is the time taken to insert a value when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?
- What is the time taken to delete the value at a given index position when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?
- What is the time taken to search for a value when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?
- What is the time taken to access a value at an index position when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?

Make a plot. For your experiments, make sure your array contains distinct

values. □

Exercise 105.5.2. Write methods to return the size of the array, boolean functions to tell you if the array is empty or full (i.e., `n` has reached the maximum capacity.) Also, throw some exceptions objects when performing invalid deletes or inserts.

We can abstract the situation and think of an `Array< T >` object as containing `T` values laid out on a straight line. The supporting methods are

```
list.insert(i, v): insert v into list at index position i
list.remove(v)   : remove one value v from list
list.remove(i)   : remove the value at index position i
list.find(v)     : return the index where v occurs (left-to-right)
                  or return -1 if not found
list[i]          : return the value at index i
list[i] = v      : change value at index i to v
list.size()      : return the size of list
list.is_empty()  : return true iff list is empty
list.is_full()   : return true iff list is full
```

Note that when you describe a container and the operations available on the container (or using OO-speak ... what the container can do) without referring to the implementation (or even the programming language used), you're describing an **abstract data type (ADT)**.

abstract data type
(ADT)

Technically speaking, the above description uses OO syntax and there are language out there which are non-OO. We would have to be even more free-form to call our description the description of an ADT. However there's no standardization ADT description just like there's no standardization of pseudocode language (!!!)

If I have to describe the above using a non-OO language say a procedural/functional language I would have to write:

```
insert(list, i, v): insert v into list at index position i
remove(list, v)   : remove one value v from list
remove(list, i)   : remove the value at index position i
find(list, v)     : return the index where v occurs (left-to-right)
                  or return -1 if not found
get(list, i)      : return the value at index i
set(list, i, v)   : change value at index i to v
size(list)        : return the size of list
is_empty(list)    : return true iff list is empty
is_full(list)     : return true iff list is full
```

Unfortunately there are others who feel the above is not free-form enough.

I'm not going to argue with the language lawyers if one should describe ADT using a pseudo-OO language or pseudo-procedural/functional language. This is how I'm going to describe our unordered list ADT ... and that's that:

```
ADT: Unordered list
list.insert(i, v): insert v into list at index position i
list.remove(v)   : remove one value v from list
list.remove(i)   : remove the value at index i from list
list.find(v)     : return the index where v occurs (left-to-right)
                  or return -1 if not found
list[i]          : return the value at index i
list[i] = v      : change value at index i to v
list.size()      : return the size of list
list.is_empty()  : return true iff list is empty
list.is_full()   : return true iff list is full
```

In general when you have a container where the value is abstractly laid out in a straight line and you have operations to insert, delete, etc on the list, you have the ADT called a **list**. Specifically, a list is either empty or among the values in the list there is one that is the head of the list and there is one that is the tail. (The head can be the tail.) Associated to every value in the list that is not a tail, there is the concept “next” value.

list

If the values in the list is not ordered in any way we say that it is an **unordered list**.

unordered list

Our `Array< T >` is an implementation of the unordered list. The language used is of course C++ and the implementation uses the C/C++ static array.

If the list is ordered, we say that the ADT is an **ordered list**. Here's the ADT. In the next section, I'll talk about the sorted array which is an example of a sorted list.

ordered list

```
ADT: Ordered list
list.insert(v)   : insert v into list
list.remove(v)   : remove one value v from list
list.remove(i)   : remove the value at index i from list
list.find(v)     : return the index where v occurs (left-to-right)
                  or return -1 if not found
list[i]          : return the value at index i
list[i] = v      : change value at index i to v
list.size()      : return the size of list
list.is_emptyp() : return true iff list is empty
list.is_full()   : return true iff list is full
```

Note that the insert method for a sorted list does not require an index value since the list will insert v in the right place since the list is ordered.

The `SortedArray< T >` is an implementation of the order list. The language used is again C++ and the static array is again used.

105.6 Static sorted array debug: static-sorted-array.tex

What if your static array is always sorted?

105.6.1 Insert

If you have to insert a value at an index position i , the runtime is still the same as the case of a static unsorted array. (Think about it. Remember the array must be sorted. Technically, since the array is sorted, you don't provide an index. We say that container is **self-organizing**. But you can also allow the user to specify an index, with the caveat that the user knows what he/she is doing.)

105.6.2 Delete

If you have to delete a value at an index position i , the runtime is still the same as the case of a static unsorted array. (Right?)

105.6.3 Search

In this case, since the array is sorted, instead of scanning the array left-to-right, you would use binary search. The worse runtime (in fact also the average) is

$$O(\lg n)$$

Of course if the value you're looking for is exactly in the middle of the array (in the sense that the first probe of binary search hits that value), then the runtime is

$$O(1)$$

Sorted array	Best	Average	Worse
Insert	$O(1)$	$O(n)$	$O(n)$
Delete	$O(1)$	$O(n)$	$O(n)$
Search	$O(1)$	$O(\lg n)$	$O(\lg n)$

Exercise 105.6.1. Write a template `SortedArray` class with the following methods:

```
void insert(T key);  
void delete(T key);  
int  search(T key); // returns index or -1 (if not found)  
T    operator[] (int index) const;  
T &  operator[] (int index);
```

Do not allow duplicate values in the `SortedArray` object. During an `insert`, if the value to be inserted is already in the object, throw a `DuplicateValue` exception object. Set the maximum capacity of your `Array` objects to 1000000. Collect timing data for the following scenarios for `SortedArray< int >` objects:

- What is the time taken to insert a value when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?
- What is the time taken to delete the value at a given index position when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?
- What is the time taken to search for a value when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?
- What is the time taken to access a value at an index position when $n = 100000, 200000, 300000, 400000, \dots, 1000000$?

Make a plot. For your experiments, make sure your array contains distinct values. □

Exercise 105.6.2. For the case of sorted and unsorted *integer* array classes, i.e., not template classes, note that a sorted integer array is an integer array. Try writing `IntArray` class and `SortedIntArray` classes using inheritance. □

105.7 Dynamic array debug: dynamic-array.tex

By a dynamic array, I mean an array in the heap or free store. Recall that a dynamic array can have variable size:

```
T * p;  
...  
p = new T[size]; // size is a not-necessarily const variable  
...  
delete [] p;  
...
```

All the runtimes are the same except when the array needs to be reallocated. For instance let's think about the insert tail. In the case of static array, if the array is not full, then the runtime is $O(1)$. If it's full, you abort mission – either you throw an exception right away or you do nothing to the array. What about the dynamic array?

If the dynamic array is not full, then the runtime is again $O(1)$. If it's full, you request for a large array. Then you need to copy your values over to the new array and perform insert tail and (don't forget!) deallocate the memory used by the smaller array. So in this case, the runtime is $O(n)$, not $O(1)$. (We're ignoring the time taken for the memory request.) Of course the memory allocation for the new array can fail – but let's ignore that.

What about delete? If the delete is delete tail? Then the runtime is $O(1)$ just like the case of static array. However if you want to make sure your dynamic array is not wasting too much memory, then you want to use a smaller array when the length of the array is less than $1/3$ of the capacity. You allocate an array of size say $2n$, copy the values over to the new array, deallocate the old array. In that case, you again have a runtime of $O(n)$.


105.8 Memory fragmentation debug: memory-fragmentation.tex

From the above, it's clear that dynamic arrays (arrays in the heap) is better than static arrays (arrays in an area of your computer memory called stack segment).


One extremely important fact about dynamic arrays is this: because of the unpredictable calls to the heap for allocation and deallocation of memory, after some time your heap becomes highly “fragmented”. What is this **memory fragmentation**?

memory
fragmentation


Suppose when you first run your program, your heap looks like this (i.e., no memory allocation yet so the whole heap is available):




Suppose your program requests for a block of 2KB of memory for pointer p0. Then your memory might look like this:



And then your program asks for another block of 1KB for `p1` and then 3KB for `p2`



And if your program deallocates the memory used by `p1`, then the heap looks like this:



After some time, your heap might look like this:



due to a huge number of memory allocations and deallocations.

Your heap memory becomes highly fragmented if you have many blocks of free memory that are not contiguous. This is a problem. Why? ...

Because you might actually have a total of 10KB of memory which is scattered all over the place and the largest *contiguous* free block of memory might only be 1KB. In that case, you cannot allocate *contiguous* free memory for a 2KB dynamic array even if you actually have a 10 KB of *total* free memory.

The fact that the memory usage of an array must be contiguous should be clear from the facts found in my chapters on pointers from CISS245: If p points to an array of values, note that the address in $p + i$ is very close to the value of $p + i + 1$. (More about low level details about memory usage in CISS360.)

We will be studying many data structures (i.e. containers of values) that will hold values in a “scattered” way to combat this very serious shortcoming of the array containers.

105.9 Accessing values in the container: index values

debug: index-values.tex

Note that in the case of an array, you have the concept of an index value of a particular value in the array. The index is of course the “position” or “location” of that value. If the array is `x`, and `i` contains a valid index value, then `x[i]` arrives at a value in `x`. You can (and *should*) think of an index value as a mechanism for referring to or reaching a value in your array.

In the case of an array, *ZZZ*

- You want to add a value to a container *at an index position*. This case potentially overwrites a value. So for instance in the case of an array, values at that index position up to the last index value of the array is moved to the right by one index position.
- You want to delete a value from the container *at an index position*.
- You want to search for a student in the container, *returning the index position if it's found*. If the student is not found, a special index value is returned to indicate failure.

Again, the index represents a position. In the case of the search, if a student is found and the index position is returned, the index value can then be used to locate the student (so as to modify satellite data, print the satellite data, etc.)

You'll see very soon that a more uniform way to access a value in a container is through a pointer. And to make the pointer more flexible, we wrap the pointer in an object. In fact sometimes, depending on what you want to do to the value the pointer is pointing to, this object can contain more than just a single pointer. Such an object (that contains a pointer – and maybe more – to access a value in a container) is called an iterator. But first let's compare the index variable and the pointer ...

105.10 Accessing values in a container: pointers

debug: pointers.tex

You usually want to traverse a container, i.e., “travel through” a container. In the case of an array, you can traverse the container by running through the index values of the array. (Like I said in the previous section, think of the index as a mechanism for accessing a value in the array.) For instance:

```
for (int i = 0; i < n; ++i)
{
    std::cout << student[i] << std::endl;
}
```

Note that when you’re doing a search, you are in fact doing a traversal. When doing a search in an array, when the target key value is found, an index value is returned (with -1 indicating search failure). Then index value is then used to access the value in the array for some work (such as printing, modification, etc.)

In any case, an index value can be used to access a value in an array.

Note that there’s another way to locate a particular student in the `student` array: by using pointers. So instead of

```
...

int find(const Student & student[], const std::string & id)
{
    ...
}

int main()
{
    ...
    std::string id;
    std::cin >> id;
    int index = find(student, id);
    std::cout << student[index] << std::endl;
    ...
}
```

we can do this:

```
...

Student * find(Student student[], std::string id)
{
    ...
}

int main()
{
    ...
    std::string id;
    std::cin >> id;
    Student * p = find(student, id);
    std::cout << (*p) << std::endl;
    ...
}
```

Compare the following traversal (using index values):

```
for (int i = 0; i < n; ++i)
{
    std::cout << student[i] << std::endl;
}
```

with this (using pointers):

```
for (const Student * p = &student[0]; p != &student[n]; ++p)
{
    std::cout << (*p) << std::endl;
}
```

Note that `&student[n]` is the address *just past* the *last* address occupied by the last student value `&student[n - 1]`. Think about the above very carefully. Drawing a picture of the computer's memory helps.

So what exactly is the difference between the two?!? And which is better? Note that the index version looks like this

```
...
    std::cout << student[i] << std::endl;
...
```

and it uses `student` and `i` whereas the pointer version

```
...
    std::cout << (*p) << std::endl;
...
```

uses `p`. Without going into assembly/machine code, it should already be clear that the index version is slower since it requires memory access to more variables (if nothing else), i.e., two variables `student` and `i` compared against one

variable `p` for the pointer version. In fact that *is* the case at the assembly code level: compilers actually translate array traversal by index into array traversal by pointer.

Another reason why the pointer version is better is because, if you think about the meaning of `student[i]`, you see that you start at the memory location of `student[0]` and go to the memory address of the i -th value in order to get to the i -value in `student`. In the case of the array, this can be done rather quickly. You will see later that there are other containers where the computation to get to the i -th value from the first value is very slow. For some containers, it's a lot easier to compute the location of the i -th value from the $(i - 1)$ -st value. In other words the method of using index values is not appropriate in some cases. I'll give you specific examples later. For now just remember that index values might not make sense for some containers.

Note that the pointer method requires us to know the beginning address of the container and the “end of address” of the container or rather the address that is just outside the address space occupied by the container's values. Nonetheless I'm going to call this “outside end of address” the end address.

105.11 Accessing values in a container: iterators

debug: iterators.tex

In general, lots of C++ STL containers come with features to access values in the containers using iterators (which are more or less pointers). For instance in the case of `std::vector` class (see CISS245), you can do the following to traverse a vector of integers `v`:

```
for (typename std::vector< int >::iterator p = v.begin();
     p != v.end(); ++p)
{
    ... do something with (*p) ...
}
```

In the above


`std::vector< int >::iterator`

is a class: it's a class inside the `std::vector` class. (Yes you can put a class inside a class.) So you would expect the `std::vector` class to look like this:

```
template < typename T >
class vector
{
public:
    class iterator
    {
    };
};
```

Objects of this class, `iterator`, are iterators to for instance `std::vector< int >` objects (and remember that iterators are like pointers).

`v.begin()` will “essentially” return a pointer to `v[0]`. OK ... it's not a pointer ... it's an iterator, i.e., an object *containing a pointer* that points to `v[0]`. `v.end()` also returns an iterator that contains a pointer that points to `&v[n]` where `n` is `v.size()`. Again, `v.end()` does not return a pointer value, but rather an *iterator* that *contains a pointer* that points to `&v[n]`. And remember that `&v[n]` is not an address of one of the values of `v` – it's an address that is *just beyond* or *just outside* of `v`. Here's a picture that represents these two iterators:



C++ STL has a container called the `std::list`. If `x` is a list (i.e., STL list) of integers, then you can do this:

```
for (typename std::list< int >::iterator p = x.begin();
     p != x.end(); ++p)
{
    ... do something with (*p) ...
}
```

A list is a kind of container that I will talk about soon: it's a doubly-linked list. There's another C++ STL container called the `std::set`. If `x` is a set of integers, then you can do this:

```
for (typename std::set< int >::iterator p = x.begin();
     p != x.end(); ++p)
{
    ... do something with (*p) ...
}
```

As you can see, the code to run through the values in the above three containers look almost exactly the same.

Get it? This is why knowing iterators is useful, because of the fact that they unify the access to values in C++ STL containers.

By the way, as I said above, the code

```
... typename std::vector< int >::iterator p ...
```

tells you that *inside* the class `std::vector< int >`, there's a class called `iterator`. In other words the `std::vector` class looks like this:

```
template < typename T >
class vector
{
public:
    class iterator
    {
    };
private:
};
```

Now you might ask ... what's this `typename` keyword thing for:

```
... typename std::vector< int >::iterator p ...
```

The `typename` tells your compiler that `std::vector< int >::iterator` is a type. Why must you do this? Because remember that from CISS245, when C++ sees `[class]::[something]`, it could mean that `[something]` is a static instance member. The default is static instance member. The `typename` tells your C++ compiler that it's a class, not a static member.

Note that for the case of `std::vector`, you do have `operator[]` so that you can do this:

```
for (size_t i = 0; i < v.size(); ++i)
{
    ... v[i] ...
}
```

i.e., you can run through a vector container using index values instead of iterators. The `size_t`, “size type”, is basically `unsigned int` and is always ≥ 0 .

The problem is that `std::list` does *not* have `operator[]`. So the index method of traversal does *not* apply to STL lists. And why doesn’t STL list have `operator[]`? Because the list container has a certain structure that makes this operator, i.e. `operator[]`, very slow if this was ever implemented. In other words, you shouldn’t ever want or think about getting the i -th value of a list using `operator[]`. That’s why it’s not included in the list library.

However arrays can go to the i -th value extremely fast. (This will be explained in CISS360.) That’s why arrays (and `std::vector` objects) have `operator[]`.

Wrapping a pointer inside an object (an iterator) is more flexible simply because there are times when the “move pointer to the next object in the container” is very complex and needs extra data. For instance suppose for some bizarre reason you have a vector `v` of 10 values and you want to access the values in this order:

`v[1], v[2], v[2], v[3], v[3], v[3], v[4], v[4], ...`

i.e., you want to process `v[i]` i times. If you do this frequently, you might want to create an iterator class that does that so that doing `++p` moves the iterator like the above. The iterator object must then remember the index value and the number of times it has processed the value at the index value. Right? Get it?? If you put all that information into the iterator object, it will simplify your (crazy) algorithm.

105.12 C++ STL: iterators and constant iterators

debug: constant-and-non-constant-iterator.tex

Implementing an iterator class is easy: refer to CISS245 if you need help. In particular, look at the notes on the `IntPtr` class. Specifically, your personal vector class should look like this:

```
template < typename T >
class vector
{
public:
    class iterator
    {
    private:
        T * q_; // points to a T value in the array that p_
                // points to
    };
private:
    T * p_; // points to an array in the free store
};
```

In the case of

```
for (typename std::vector< int >::iterator p = v.begin();
     p != v.end(); ++p)
{
    ... do something with (*p) ...
}
```

The iterator `p` can modify the value that it points to:

```
for (std::vector< int >::iterator p = v.begin();
     p != v.end(); ++p)
{
    (*p) = 0;
}
```

In case you do *not* want that to happen, or that you cannot do this (for instance `v` is a constant and therefore you cannot change the values in `v`), you do this:

```
for (typename std::vector< int >::const_iterator p = v.begin();
     p != v.end(); ++p)
{
    ... only read access to (*p), not write access ...
}
```

i.e., `p` is a **constant iterator**.

constant iterator

Exercise 105.12.1. Rewrite the following using iterators:

```
template< typename T >
std::ostream & operator<<(std::ostream & cout,
                        const std::vector< T > & v)
{
    std::string delim = "";
    cout << '{';
    for (size_t i = 0; i < v.size(); ++i)
    {
        cout << delim << v[i];
        delim = ", ";
    }
    cout << '}';
    return cout;
}
```

Test it.



Exercise 105.12.2. Add a `const_iterator` class inside your `vector` class. Note that if `p` is a constant iterator to (say) a C++ STL vector object, then `*p` refers to a value in that vector object. `*p` has read access but not write access, i.e.,:

```
int a = *p; // OK
*p = 42;    // BAD!!!
```

This means that in your constant iterator class inside your vector class you have

```
template < typename T >
class vector
{
    class const_iterator
    {
    public:
        T operator*() const
        { ... }
    };
};
```

It's true that the actual value in the vector is not changed since you're returning a copy of that value. But this might be misleading to the person using your constant iterator class (including yourself if you're not careful). It's better to do this:

```
template < typename T >
class vector
{
    class const_iterator
    {
    public:
        const T & operator*() const
        { ... }
    };
};
```

Why? (You ought to know.)



You *should* write a complete template vector class that has supporting iterators (constant and nonconstant).

105.13 C++ STL: `std::vector` operations/methods using iterators debug:

vector-methods-with-iterators.tex

Here are some examples of `std::vector` methods that uses iterators. Let `v` be a `std::vector< T >` object. Let `p`, `q`, `r` be iterators.

<code>v.begin()</code>	returns iterator pointing to first value of <code>v</code>	
<code>v.end()</code>	returns iterator pointing to just beyond last value of <code>v</code>	
<code>v.insert(p, x)</code>	insert <code>x</code> at <code>p</code>	<code>std::vector::insert</code>
<code>v.insert(p, q, r)</code>	insert values at <code>q</code> , ..., <code>r - 1</code> into <code>v</code> at <code>p</code>	
<code>v.erase(p)</code>	erase <code>v</code> at <code>p</code>	<code>std::vector::erase</code>
<code>v.erase(p, q)</code>	erase <code>v</code> at <code>p</code> , ..., <code>q-1</code>	
<code>std::find(p, q, t)</code>	find <code>t</code> at <code>p</code> , ..., <code>q-1</code> . If not found, <code>v.end()</code> is returned.	<code>std::find</code>
<code>std::sort(p, q)</code>	sort values at <code>p</code> , ..., <code>q-1</code>	<code>std::sort</code>
<code>std::distance(p, q)</code>	number of values at <code>p</code> , ..., <code>q-1</code>	<code>std::distance</code>

Make sure you study the code below carefully.

```
#include <iostream>
#include <string>
#include <vector>
#include <algorithm> // for std::find, std::sort, std::distance

template< typename T >
std::ostream & operator<<(std::ostream & cout,
                        const std::vector< T > & x)
{
    std::string delim = "";
    cout << '{';
    for (typename std::vector< T >::const_iterator p = x.begin();
         p != x.end(); ++p)
    {
        cout << delim << (*p);
        delim = ", ";
    }
    cout << '}';
    return cout;
}

int main()
{
    std::vector< int > v = {13, 4, 5, 10, 57, 23, 52, 12, 7};
    std::cout << "v: " << v << '\n';
}
```

```

typename std::vector< int >::iterator p = v.begin();
std::cout << (*p) << ' ';
++p;
std::cout << (*p) << ' ';
p++;
std::cout << (*p) << ' ';
p += 2;
std::cout << (*p) << ' ';
p -= 3;
std::cout << (*p) << '\n';

p = v.begin() + 3;
v.insert(p, -1); // insert value
std::cout << "v: " << v << '\n';

std::vector< int > u = {-1, -2, -3, -4, -5, -6};
p = v.begin() + 1;
v.insert(p, u.begin() + 1, u.end() - 2); // insert range
std::cout << "v: " << v << '\n';

p = v.begin();
v.erase(p + 5); // delete value
std::cout << "v: " << v << '\n';

p = v.begin();
v.erase(p, p + 3); // delete range
std::cout << "v: " << v << '\n';

std::sort(v.begin(), v.end() - 2); // sort range
std::cout << "v: " << v << '\n';

p = std::find(v.begin(), v.end(), 10); // find in a range
if (p != v.end()) // (success case)
{
    std::cout << "10 found\n";
}
p = std::find(v.begin(), v.end(), 9999); // find in a range
if (p == v.end()) // (failure case)
{
    std::cout << "9999 not found\n";
}

std::cout << "v: " << v << '\n';
p = v.begin() + 2;

```

```
typename std::vector< int >::iterator q = v.end() - 1;
std::cout << (*p) << ' ' << (*q) << ' '
          << std::distance(p, q) // number of values in range
          << '\n';              // (value at q not included)

std::cout << std::vector< int >{2, 3, 5, 7, 11, 13} << '\n';

return 0;
}
```

Here's the output:

```
[student@localhost containers] g++ main.cpp; ./a.out
v: {13, 4, 5, 10, 57, 23, 52, 12, 7}
13 4 5 57 4
v: {13, 4, 5, -1, 10, 57, 23, 52, 12, 7}
v: {13, -2, -3, -4, 4, 5, -1, 10, 57, 23, 52, 12, 7}
v: {13, -2, -3, -4, 4, -1, 10, 57, 23, 52, 12, 7}
v: {-4, 4, -1, 10, 57, 23, 52, 12, 7}
v: {-4, -1, 4, 10, 23, 52, 57, 12, 7}
10 found
9999 not found
v: {-4, -1, 4, 10, 23, 52, 57, 12, 7}
4 7 6
{2, 3, 5, 7, 11, 13}
```

In the above example, note that the initializer list for `std::vector` initialization (i.e., constructor)

```
std::vector< int > u = {-1, -2, -3, -4, -5, -6};
std::cout << std::vector< int >{2, 3, 5, 7, 11, 13} << '\n';
```

is available in g++ (as per the C++11 standard) of our current fedora virtual machine. (For more information, see section on initializer lists.)

105.14 C++: Automatic type deduction debug:

type-deduction.tex

[NOTE: See also CISS362 notes.]

Later versions of g++ (C++11 and later) support automatic type deductions and range-based for-loops. Here's an easy example on automatic type deduction:

auto

```
#include <iostream>
#include <set>

int main()
{
    auto i = 42;
    auto x = 3.14;
    auto j = i;
    auto & k = i;
    std::cout << i << ' ' << x << ' ' << j << ' ' << k << '\n';
    i = -1;
    std::cout << i << ' ' << x << ' ' << j << ' ' << k << '\n';

    return 0;
}
```

```
[student@localhost containers] g++ main.cpp; ./a.out
42 3.14 42 42
-1 3.14 42 -1
```

In this case g++ figured out that `i` should have type `int` and `x` should be type `double`.

(To understand how type deductions work, see CISS445 where I'll talk about type inferencing for the OCAML language.)

Just because you can use `auto`, it does *not* mean you can forget about what is the actual type you want. Using something blindly is dangerous. Frequently going back to explicit code is helpful. You should learn the concept of C++ type deduction and ranged-based for-loop. You might want to use it to quickly write code for assignments. But once you are done with the assignment, you should replace the type deduction code and range-based for-loops with the explicit version so that you don't forget what is the C++ code generated. Of course you can use it in your personal projects.

105.15 C++: Range-based loops debug: range-based-for-loops.tex

Here's an example of **range-based for-loop** where `x` is a `std::vector< int >` range-based for-loop object:

```
for (int i: x)
{
    std::cout << i << '\n';
}
```

This is frequently used together with automatic type deduction: auto

```
for (auto i: x)
{
    std::cout << i << '\n';
}
```

Here's an example:

```
#include <iostream>
#include <string>
#include <vector>

template< typename T >
std::ostream & operator<<(std::ostream & cout,
                        const std::vector< T > & x)
{
    std::string delim = "";
    cout << '{';
    for (auto i: x)
    {
        cout << delim << i;
        delim = ", ";
    }
    cout << '}';
    return cout;
}

int main()
{
    int w[] = {1,2,3};
    for (int i: w)
    {
        std::cout << i << ' ';
    }

    std::vector< int > v = {2, 3, 5, 7, 11, 13};
    std::cout << v << '\n';

    auto p = v.begin();
    ++p;
```



```

    std::cout << (*p) << '\n';
    auto q = v.end();
    --q;
    v.erase(p, q);
    std::cout << v << '\n';

    return 0;
}

```

```

[student@localhost containers] g++ main.cpp; ./a.out
1 2 3 {2, 3, 5, 7, 11, 13}
3
{2, 13}

```

If you want to modify a value in a container, you should use this in the declaration of `v`:

```

for (auto & v: x)
{
    v = 0;
}

```

so that `v` is a reference to a value in `x`. Of course you cannot change the values of `x` if `x` is constant. So in the above I'm assuming `x` is not constant. That's because for

```

for (auto v: x)
{
    v = 0;
}

```

a separate local copy of a value in `x` is created for `v` and changing `v` to 0 will not change the corresponding value in `x`. And of course it's slow especially if the values in container `x` are complex and time consuming to construct. There are some cases where you have to do *this* instead

```

for (auto && v: x)
{
    v = 0;
}

```

if you want to change the values in `x`. So I suggest you use this `&&` method anyway if you want to change the values in `x`.

Usually you should make `v` a reference to a value in `x` and not make copies all the time. If `x` is constant and your loop does not change the values in `x`, you can do this when the values in `x` are complex (example: objects):

```
for (const auto & v: x)
{
    ...
}
```

Note that if your `x` is constant and you do

```
for (auto & v: x)
{
    ...
}
```

the type deduction will be smart enough to make the `v` a constant reference. There are many other variations.

In summary if you do not want to change the values in `x` and the values are simple and you don't mind copying the values of `x` to `v`, do this:

```
for (auto v: x)
{
    ...
}
```

If you do not want to change the values in `x` and you want `v` to reference values in `x`, you should do this:

```
for (const auto & v: x)
{
    ...
}
```

(the `const` is redundant if `x` is constant). If `x` is not constant and you want to change the values in `x`, you can do

```
for (auto & v: x)
{
    ...
}
```

or

```
for (auto && v: x)
{
    ...
}
```

where the second version is probably better.

The following is an example involving `std::vector`:

```
#include <iostream>
#include <string>
#include <vector>
```

```
template< typename T >
std::ostream & operator<<(std::ostream & cout,
                        const std::vector< T > & x)
{
    std::string delim = "";
    cout << '{';
    for (auto & v: x)
    {
        cout << delim << v;
        delim = ", ";
    }
    cout << '}';
    return cout;
}

int main()
{
    std::vector< int > v = {2, 3, 5, 7, 11, 13};
    std::cout << v << '\n';

    for (auto x: v)
    {
        x = 0;
    }
    std::cout << v << '\n';

    for (auto & x: v)
    {
        x = 1;
    }
    std::cout << v << '\n';

    for (auto && x: v)
    {
        x = 2;
    }
    std::cout << v << '\n';

    return 0;
}
```

```
[student@localhost containers] g++ main.cpp; ./a.out
{2, 3, 5, 7, 11, 13}
{2, 3, 5, 7, 11, 13}
{1, 1, 1, 1, 1, 1}
{2, 2, 2, 2, 2, 2}
```

And again these are not the only possibilities. And if the compiler does not

deduce the right type or your code does not allow a suitable type deduction, I suggest you go back to basics and type in the actual type you really want.

In some STL containers, the values cannot be changed. For instance the values in a `std::set` object cannot be changed – they are immutable. Of course you know that values in a `std::vector` can be changed if it's not constant.

Again be aware of the above syntax. Try out the above examples. But avoid using it in assignments until you really know your iterators without using shortcuts like `auto` and range-based for-loops.

105.16 C++ STL: `std::pair` and `std::tuple` debug:

products-and-tuples.tex

C++ provides two STL classes for tuples: `std::pair` for 2-tuples and `std::tuple` for general k -tuples.

Each `std::pair` object is made up of two values which can be of different types. If `x` is a `std::pair` object, then the two values are `x.first` and `x.second`. In other words, `std::pair` is basically just this: first
second

```
template < typename S, typename T >
class pair
{
public:
    ...
    S first;
    T second;
};
```

(Actually `std::pair` is a `struct`.)

Each `std::tuple` object is made up of two or more values which can be of different types. The number of value is fixed. If `x` is a `std::tuple` object, then the values are `std::get<0>(x)`, `std::get<1>(x)`, `std::get<2>(x)`, etc. std::get

Run the following example code:

```
#include <iostream>
#include <set>
#include <utility> // for std::pair
#include <tuple>

template< typename S, typename T >
std::ostream & operator<<(std::ostream & cout,
                        const std::pair< S, T > & x)
{
    cout << '(' << x.first << ", " << x.second << ')';
    return cout;
}

template< typename S, typename T, typename U >
std::ostream & operator<<(std::ostream & cout,
                        const std::tuple< S, T, U > & x)
{
    cout << '('
        << std::get<0>(x) << ", "
```

```

        << std::get<1>(x) << ", "
        << std::get<2>(x) << ' ';
    return cout;
}

int main()
{
    std::pair< int, double > u = {2, 3.14159};
    std::cout << u << '\n';
    u.first = 1;
    u.second = 2.718281;
    std::cout << u << '\n';

    std::tuple< int, double, char > v = {-2, 0.01, 'A'};
    std::cout << v << '\n';
    std::get<0>(v) = -1;
    std::get<1>(v) = -4.2;
    std::get<2>(v) = 'B';
    std::cout << v << '\n';

    return 0;
}

```

```

[student@localhost containers] g++ main.cpp; ./a.out
(2, 3.14159)
(1, 2.71828)
(-2, 0.01, A)
(-1, -4.2, B)

```

Of course a `std::pair` is just a special case of `std::tuple`.

It's obvious that the length of the tuple in the above example is fixed at 3. If you want to have a tuple of arbitrary length you should not use `std::tuple`. If the values in your tuple have the same type but you want the length to be arbitrary, then you should use a `std::vector`. Of course a `std::vector` is homogeneous: the type of the values must be the same. If you want to have a `std::vector` of “different types” of objects, say from class A, B, C, you can create a parent class P for A, B, C, and use `std::vector< P * >` and polymorphism (see CISS245).

105.17 C++: typedefs debug: typedefs.tex

Besides the fact that automatic type deductions and range-based for-loops makes container code easier to write, **typedefs** help too. You have already seen **typedefs** from CISS240, CISS245. (You should check my CISS240, CISS245 notes on typedefs.)

Suppose you want to work with a pair of doubles and also an array of pairs of doubles. You can do this:

```
#include <iostream>
#include <vector>

std::ostream & operator<<(std::ostream & cout,
                        const std::pair< double, double > & x)
{
    cout << '(' << x.first << ", " << x.second << ')';
    return cout;
}

std::ostream & operator<<(std::ostream & cout,
                        const std::vector< std::pair< double, double > > & v)
{
    cout << '{';
    std::string delim = "";
    for (auto & x: v)
    {
        cout << delim << x;
        delim = ", ";
    }
    cout << '}';
    return cout;
}

int main()
{
    std::pair< double, double > pair;
    std::vector< std::pair< double, double > > pairs;

    return 0;
}
```

Or you can use typedefs:

```
#include <iostream>
#include <vector>

typedef std::pair< double, double > twodouble;
```

```
typedef std::vector< std::pair< double, double > > twodoubles;

std::ostream & operator<<(std::ostream & cout, twodouble & x)
{
    cout << '(' << x.first << ", " << x.second << ')';
    return cout;
}

std::ostream & operator<<(std::ostream & cout, twodoubles & v)
{
    cout << '{';
    std::string delim = "";
    for (auto & x: v)
    {
        cout << delim << x;
        delim = ", ";
    }
    cout << '}';
    return cout;
}

int main()
{
    twodouble pair;
    twodoubles pairs;

    return 0;
}
```

Clearly the typedefs help!

105.18 C++: initializer lists debug: initializer-lists.tex

The “initializer list” refers to the array initializer from C11:

```
int x[] = {2, 3, 5};
```

and not the constructor initializer list from C11:

```
class C
{
public:
    C()
        : x_(2), y_(3), z(5)
    {}
};
```

Note that the original notation for (array) initializer list can only be used during declaration.

In the newer C++ (since C++11), the **initializer list** notation can be used when you work with C++ STL containers and it can be used for assignments.

initializer list

Before the introduction of initializer lists for STL template classes, it was very tedious to initialize an STL container. You can first initialize it to an empty container and then you put values into the container one value at a time:

```
std::vector< int > v;
v.push_back(2);
v.push_back(3);
v.push_back(5);
v.push_back(7);
```

You can also do this:

```
int x[] = {2, 3, 5, 7};
std::vector< int > v(x, x + 4);
```

But you do have to create an array of the values and the unnecessary name of `x`. Both cases are tedious.

With the new initializer lists for STL template classes, you can do the following during constructor call with initialization:

```
std::vector< int > v = {2, 3, 5, 7};
```

or

```
std::vector< int > v {2, 3, 5, 7};
```

You can also use the initializer list notation with the assignment operator

```
std::vector< int > v;
v = {2, 3, 5, 7};
```

Run and study the following:

```
#include <iostream>
#include <vector>

typedef std::pair< double, double > twodouble;
typedef std::vector< twodouble > twodoubles;

std::ostream & operator<<(std::ostream & cout, const twodouble & x)
{
    cout << '(' << x.first << ", " << x.second << ')';
    return cout;
}

std::ostream & operator<<(std::ostream & cout, const twodoubles & v)
{
    cout << '{';
    std::string delim = "";
    for (auto & x: v)
    {
        cout << delim << x;
        delim = ", ";
    }
    cout << '}';
    return cout;
}

int main()
{
    twodoubles v = {{0.0, 0.1}, {0.2, 0.3}, {0.4, 0.5}};
    std::cout << v << '\n';

    v = {{1.1, 2.2}};
    std::cout << v << '\n';

    std::cout << (twodoubles {{0.0, 0.1}, {0.2, 0.3}, {0.4, 0.5}}) << '\n';

    return 0;
}
```

Here's the output:

```
[student@localhost containers] g++ main.cpp; ./a.out
{(0, 0.1), (0.2, 0.3), (0.4, 0.5)}
{(1.1, 2.2)}
{(0, 0.1), (0.2, 0.3), (0.4, 0.5)}
```

For

```
twodoubles v = {{0.0, 0.1}, {0.2, 0.3}, {0.4, 0.5}};
```

without using typedefs, it would become

```
std::vector< std::pair< double, double > > v  
    = {{0.0, 0.1}, {0.2, 0.3}, {0.4, 0.5}};
```

Without the inner initializer list, it becomes

```
std::vector< std::pair< double, double > > v  
    = {std::pair< double, double >(0.0, 0.1),  
      std::pair< double, double >(0.2, 0.3),  
      std::pair< double, double >(0.4, 0.5)};
```

i.e., initializer lists have to be replaced by explicit constructor calls. And without the outer initializer list it becomes

```
std::vector< std::pair< double, double > > v;  
v.push_back(std::pair< double, double >(0.0, 0.1));  
v.push_back(std::pair< double, double >(0.2, 0.3));  
v.push_back(std::pair< double, double >(0.4, 0.5));
```

which is truly horrific.

By the way in the above, just like the case for `std::vector` I did this:

```
twodoubles v = {{0.0, 0.1}, {0.2, 0.3}, {0.4, 0.5}};
```

I can also do this:

```
twodoubles v {0.0, 0.1}, {0.2, 0.3}, {0.4, 0.5};
```

Finally the initializer notation can be used in a constructor call to create an object value without assigning it to a name:

```
std::cout << (twodoubles {0.0, 0.1}, {0.2, 0.3}, {0.4, 0.5}) << '\n';
```

In the same way you can do

```
std::cout << (std::vector< int > {2, 3, 5}) << '\n';
```

which is very convenient for short tests:

```
std::cout << somesort(std::vector< int > {2, 5, 3}) << '\n';
```

105.19 C++ STL `std::vector` exercises debug: exercises.tex

You know what you should: practice all the new C++ concepts (iterators and initializer list) and redo as many array algorithms as you can.

Exercise 105.19.1. Implement a linear search function on a `std::vector< T >` object `v` using the concept of iterator. (This is just the `std::find`.) The return value is an iterator (not an index). If the target value is not found, return `v.end()`.

Exercise 105.19.2. Implement a bubblesort function on a `std::vector< T >` object using iterators.

Exercise 105.19.3. Implement a binary search function on a `std::vector< T >` object `v` using iterators. If the target value is not found, return `v.end()`. (Hint: You'll need `std::distance`.)

Exercise 105.19.4.

1. Create a function `randvec(n, a, b)` that returns a vector of size `n` of random integers selected from `[a, b)`.
2. Create a function `randvec(n, p, q)` that returns a vector of size `n` of random values selected from `*p, *(p+1), ..., *(q-1)` (note: `*q` is not included).

Exercise 105.19.5. Write a function `merge(u, v)` that accepts two sorted vectors and return a sorted vector from the values in `u` and `v`. Use the obvious algorithm that runs in $O(n)$. (Hint: scan both vectors at the same time and choose the smaller value from the two to be placed in the new vector.)

Exercise 105.19.6. Implement a bubblesort function on a `std::vector` of values of the form `(a, b)` where `a` is a student name (of `std::string` type) and `b` is a student id (of `int` type). Sort in ascending order based on the student id. You must use `std::pair`.

Exercise 105.19.7. Let `v` be a `std::vector` of points in 2D space, i.e., each value of `v` is of the form `(x, y)` where `x` and `y` are doubles. The `std::pair` is used. You want to sort `v` based on the distance from `(0, 0)`. The distance of

(x, y) from $(0, 0)$ is given by the distance formula (or Pythagorus theorem). Since you will need to compute the distance of (x, y) from the origin several times, you have decided to do this: you create a new `std::vector`, called `v1`, of (x, y, d) where d is the distance of (x, y) from the origin. After your bubblesort, you copy the contents of `v1` to `v` (of course omitting the d value). Test your program.

Exercise 105.19.8. In the above exercise, for each `std::pair< double > {x, y}` of `v`, you create a `std::tuple< tuple > {x, y, d}` in `v1`. Now do this: create `std::pair< std::pair< double >, double > {{x, y}, d}` for your values in `v1`. (Read the above carefully!) Redo the above exercise.

Exercise 105.19.9. Implement your own `vector` class (see CISS245 assignment) and include an `iterator` class and `const_iterator` class inside you `vector` class and should support dereferencing operator. (See for instance the `IntPtr` class example in the CISS245 notes.) Your `vector` class should includes a `begin()` and an `end()` method. □

Index

`&&`, [5046](#)

abstract data type (ADT), [5022](#)

`auto`, [5044](#), [5045](#)

constant iterator, [5038](#)

`first`, [5050](#)

`head`, [5015](#)

initializer list, [5054](#)

`key`, [5011](#)

`list`, [5023](#)

memory fragmentation, [5028](#)

ordered list, [5023](#)

range-based for-loop, [5045](#)

`satellite`, [5011](#)

`second`, [5050](#)

self-organizing, [5025](#)

`size_t`, [5037](#)

`std::distance`, [5041](#)

`std::find`, [5041](#)

`std::get`, [5050](#)

`std::pair`, [5050](#)

`std::sort`, [5041](#)

`std::tuple`, [5050](#)

`std::vector::erase`, [5041](#)

`std::vector::insert`, [5041](#)

`tail`, [5016](#)

`traverse`, [5002](#)

unordered list, [5023](#)